Run	**Train data (#sample)**	**CG R@1**	**WIT R@1**	**Bloom R@1**	**Mammoth R@1**	**ImageNet**	**R@1-Avg**
SEA-CLIP-Tiny	CC12M + CG-OE-filt + WIT-hf + Bloom + Mammoth-VL-SEA *(12.72M)*	29.3	21.8	23.3	15.3	36.4	**17.4**
			Dataset Ablation				
Only SEA datasets	CG-OE-filt + WIT-hf-base + Bloom *(1.21M)*	**63.7**	**54.6**	**24.9**	0.4	5.6	9.4
Only CC12M	CC12M (10.97M)	0.1	1.9	1.2	0.5	**40.0**	6.5
Only WIT	WIT-hf-base (487K)	0.2	27.7	2.2	0.1	1.3	6.6
Only Bloom	Bloom-only (21K)	0.0	0.0	1.0	0.0	0.1	4.7
Only CG	CG-only (703K)	46.5	0.2	0.6	0.0	0.2	5.1
Only Mammoth	SEA-Mammoth (540K)	—	—	—	—	—	—
CC12M + SEA	CC12M + CG-OE-filt + WIT-hf + Bloom *(12.18M)*	33.3	25.2	16.1	**0.8**	37.4	**14.2**
			KD Loss Ablation				
SEA-CLIP-Tiny w/o KD losses	CC12M + CG-OE-filt + WIT-hf + ML *(12.33M)*	47.4	41.5	28.8	—	—	—
SEA-CLIP-Tiny (CKD only)	CC12M + CG-OE-filt + WIT-hf + Bloom + Mammoth-VL-SEA *(12.87M)*	35.6	28.0	23.2	—	—	—
Teacher	—	*12.9*	*32.2*	*9.6*	*16.6*	*71.1*	*46.4*

Table 1: The dataset ablation study. All models are trained on the same training pipeline and hyperparameters. SEA-CLIP-Tiny is our full data blend (CC12M + SEA + Mammoth-VL-SEA). The dataset ablation swaps in SEA blend, CC12M, a single SEA source (WIT, Bloom, CG, or Mammoth), or CC12M+SEA as the only training data. We use `R@1-Avg` from Total Average@1 in Table **??**.