Run	**Train data (#sample)**	**CG R@1**	**WIT R@1**	**Bloom R@1**	**ImageNet**	**R@1-Avg**
SEA-CLIP-Tiny-c1 (SEA blend)	CG-OE-filt + WIT-hf-base + Bloom *(1.21M)*	**63.7**	**54.6**	**24.9**	5.6	9.4
SEA-CLIP-Tiny-c2 (CC12M)	CC12M (10.97M)	0.1	1.9	1.2	**40.0**	6.5
SEA-CLIP-Tiny-c3 (CC12M + SEA)	CC12M + CG-OE-filt + WIT-hf + Bloom *(12.18M)*	33.3	25.2	—	37.4	**14.2**
SEA-CLIP-Tiny-WIT	WIT-hf-base (487K)	0.2	27.7	2.2	1.3	6.6
SEA-CLIP-Tiny-Bloom	Bloom-only (21K)	0.0	0.0	—	0.1	4.7
SEA-CLIP-Tiny-CG	CG-only (703K)	46.5	0.2	—	0.2	5.1
SEA-CLIP-Tiny w/o KD losses	CC12M + CG-OE-filt + WIT-hf + ML *(12.33M)*	47.4	41.5	28.8	30.0	12.0
SEA-CLIP-Tiny + ML + Mammoth-VL-SEA	CC12M + CG-OE-filt + WIT-hf + ML + Mammoth-VL-SEA *(12.87M)*	33.5	31.3	19.5	35.9	18.5
SEA-CLIP-Tiny (CKD only)	CC12M + CG-OE-filt + WIT-hf + Bloom + Mammoth-VL-SEA *(12.87M)*	35.6	28.0	23.2	31.9	15.2
Teacher	—	*12.9*	*32.2*	*9.6*	*71.1*	*46.4*

Table 1: The dataset ablation study. All models are trained on the same training pipeline and hyperparameters. We ran the ablation on SEA-CLIP-Tiny using 3 choices: (c1) SEA blend, (c2) CC12M, and (c3) CC12M+SEA. We also tested on a single source of the SEA dataset, i.e., only WIT, Bloom, or CulturalGround (CG). We use `R@1-Avg` from Total Average@1 in Table **??**.